

SourceMind Test Document

This is a sample PDF file used for testing the ingestion pipeline. It contains multiple paragraphs of text that will be extracted, chunked, embedded, and stored in ChromaDB.

The PDF parser uses PyMuPDF to read the raw bytes and extract text content from each page of the document.